

# Systems Programming Cheat Sheet

## Multiprogramming

### Mutex locks

A *mutex* is a lock that we set before using a shared resource and release after using it. When the lock is set, *no other thread can access the locked region of code*. So we see that even if thread 2 is scheduled while thread 1 was not done accessing the shared resource and the code is locked by thread 1 using mutexes then thread 2 cannot even access that region of code. So this *ensures a synchronized access of shared resources in the code*.

```
pthread_t tid[2];
int counter;
pthread_mutex_t lock;
void* doSomething(void *arg) {
    pthread_mutex_lock(&lock);
    unsigned long i = 0; counter += 1;
    printf("\n Job %d started\n", counter);
    for(i=0; i<(0xFFFFFFFF);i++);
    printf("\n Job %d finished\n", counter);
    pthread_mutex_unlock(&lock);
    return NULL;
} int main(void) {
    int i = 0; int err;
    if (pthread_mutex_init(&lock, NULL) != 0) {
        printf("\n mutex init failed\n");
        return 1; }
    while(i < 2) {
        err = pthread_create(&(tid[i]), NULL, &
        doSomething, NULL);
        if (err != 0)
            printf("\ncan't create thread :[%s]",
            strerror(err));

        i++; }
    pthread_join(tid[0], NULL);
    pthread_join(tid[1], NULL);
    pthread_mutex_destroy(&lock);
    return 0; }
```

A mutex contains three things: A flag which is a 0 or a 1 (locked or unlocked), Owner which is a thread ID, Queue which holds suspended threads. It can only be unlocked by what has locked it.

### Signal handling

A **signal** is a condition that may be reported during program execution, and can be ignored, handled specially, or, as is the default, used to terminate the program.

```
FILE *temp_file;
void leave(int sig) {
    fprintf(temp_file, "\nInterrupted...\n");
    fclose(temp_file);
    exit(sig); }
main() {
    (void) signal(SIGINT, leave);
    temp_file = fopen("tmp", "w");
    for(;;) {
        printf("Ready...\n");
        (void) getchar();
    }
    exit(EXIT_SUCCESS); }
```

Forking returns the ID of the new process. The child process gets its very own process ID. One process calls fork, and if the return is negative one, the fork failed, if it's zero, you're in the child, and if it's anything larger than 0, you have the process ID of the child.

### Fork

The `fork()` system call will spawn a new child process which is an identical process to the parent except that has a new system process ID. The process is copied in memory from the parent and a new process structure is assigned by the kernel. The return value of the function is which discriminates the two threads of execution. A zero is returned by the fork function in the child's process.

The environment, resource limits, umask, controlling terminal, current working directory, root directory, signal masks and other process resources are also duplicated from the parent in the forked child process.

```
char * message; int n;
pid_t pid = fork();
switch (pid) {
    case -1:
        perror("fork failed\n");
        exit(1);
    case 0:
        message = "This is the child";
        n = 5;
        break;
    default:
        message = "This is the parent";
        n = 3;
        break; }
for(; n > 0; n--) {
    puts(message);
    sleep(1); }
```

### Exec

The `exec()` family of functions will initiate a program from within a program. They are also various front-end functions to `execve()`. The functions return an integer error code.

The function call `execl()` initiates a new program in the same environment in which it is operating. An executable (with fully qualified path. i.e. `/bin/ls`) and arguments are passed to the function. Note that `arg0` is the command/file name to execute.

Use `exec` to make a child process execute a new program after it has been forked. It usually does not return, with `-1` on failure. It will fail when: too big, ACCESS, get into a loop, the name is too long, or the executable does not exist. The new process keeps the set of blocked signals, pending signals, timers, and any open file descriptors. Doing an `exec` does not change the relationship between a parent and child process. Caught signals return default values.

```
execl("/bin/ls", "/bin/ls", "-r", "-t", "-l", (
    char *) 0);
```

### Wait

`wait()`: Blocks calling process *until the child process terminates*. If child process has already terminated, the `wait()` call returns immediately. if the calling process has multiple child processes, the function returns when one returns.  
`waitpid()`: Options available to block calling process for a particular child process not the first one.

### Shell scripting

### Shared memory

Possible for multiple processes to share segments of memory, not full memory. Beneficial for better speed. Access by multiple processes. Still need mutexes and semaphore for synchronizations. Lifetime is independent of any process. All shared memory segments exist in a single list. Identified by a key, generated by the metainfo of files and directories. Processes attach and detach to shared memory segments.

### Pointers to functions

```
void func(int);
main() {
    void (*fp)(int);
    fp = func;
    (*fp)(1);
    fp(2);
    exit(EXIT_SUCCESS); }
void func(int arg){
    printf("%d\n", arg);
}
```

### Condition variables

*Condition variables* provide yet another way for threads to synchronize. While mutexes implement synchronization by controlling thread access to data, condition variables *allow threads to synchronize based upon the actual value of data*.

Without condition variables, the programmer would need to have threads continually polling (possibly in a critical section), to check if the condition is met. This can be very resource consuming since the thread would be continuously busy in this activity. *A condition variable is a way to achieve the same goal without polling*. A condition variable is *always used in conjunction with a mutex lock*.

Used for waiting. Contains 3 basic operations: wait, signal, broadcast. Wait atomically does unlock and block. Signal atomically does re-locks a thread has it start again.

### Deadlock

Four conditions for deadlock: **mutual exclusion** (At least one resource must be held in a non-sharable mode. Only one process can use the resource at any given instant of time.), **circular wait** (A process must be waiting for a resource which is being held by another process, which in turn is waiting for the first process to release the resource.), **no pre-emption** (The operating system must not de-allocate resources once they have been allocated; they must be released by the holding process voluntarily.), **hold and wait** (A process is currently holding at least one resource and requesting additional resources which are being held by other processes.).

## Semaphores

A *semaphore* is a special type of variable that can be incremented or decremented, but *crucial access to the variable is guaranteed to be atomic*, even in a multi-threaded program. If two or more

threads in a program attempt to change the value of a semaphore, the system guarantees that all the operations will in fact *take place in sequence*.

Contains a counter, not an owner. The counter is anything that is

non-negative and can be posted by anything. Wait has two checks: if counter is not zero, decrement and return immediately, but if it is zero then suspend the calling thread and put it in the semaphore queue.