

Systems Programming Cheat Sheet

Mutex locks

A *mutex* is a lock that we set before using a shared resource and release after using it. When the lock is set, *no other thread can access the locked region of code*. So we see that even if thread 2 is scheduled while thread 1 was not done accessing the shared resource and the code is locked by thread 1 using mutexes then thread 2 cannot even access that region of code. So this *ensures a synchronized access of shared resources in the code*.

```
pthread_t tid[2];
int counter;
pthread_mutex_t lock;
void* doSomething(void *arg) {
    pthread_mutex_lock(&lock);
    unsigned long i = 0; counter += 1;
    printf("\n Job %d started\n", counter);
    for(i=0; i<(0xFFFFFFFF);i++);
    printf("\n Job %d finished\n", counter);
    pthread_mutex_unlock(&lock);
    return NULL;
} int main(void) {
    int i = 0; int err;
    if (pthread_mutex_init(&lock, NULL) != 0) {
        printf("\n mutex init failed\n");
        return 1; }
    while(i < 2) {
        err = pthread_create(&(tid[i]), NULL, &
        doSomething, NULL);
        if (err != 0)
            printf("\ncan't create thread :[%s]",
            strerror(err));

        i++; }
    pthread_join(tid[0], NULL);
    pthread_join(tid[1], NULL);
    pthread_mutex_destroy(&lock);
    return 0; }
```

A mutex contains three things: A flag which is a 0 or a 1 (locked or unlocked), Owner which is a thread ID, Queue which holds suspended threads. It can only be unlocked by what has locked it.

Signal handling

A **signal** is a condition that may be reported during program execution, and can be ignored, handled specially, or, as is the default, used to terminate the program.

```
FILE *temp_file;
void leave(int sig) {
    fprintf(temp_file, "\nInterrupted...\n");
    fclose(temp_file);
    exit(sig); }
main() {
    (void) signal(SIGINT, leave);
    temp_file = fopen("tmp", "w");
    for(;;) {
        printf("Ready...\n");
        (void) getchar();
    }
    exit(EXIT_SUCCESS); }
```

Threads

Threads are cheaper than new processes. All threads in the process share the same address space. Every thread has its own: set of registers, call stack, errno, and threadID. Advantages: Programs can keep going and doing other things even if a thread is blocked by IO and shared resources.

```
void *threadFunc(void *arg) {
    char *str; int i = 0;
    str=(char*)arg;
    while(i < 110 ) {
        usleep(1);
        printf("threadFunc says: %s\n",str);
        ++i; }
    return NULL; }
int main(void) {
    pthread_t pth; int i = 0;
    pthread_create(&pth, NULL, threadFunc, "foo");
    while(i < 100) {
        usleep(1);
        printf("main is running...\n");
        ++i; }
    printf("main waiting for thread to terminate
    ...\n");
    pthread_join(pth, NULL);
    return 0; }
```

Multiprogramming

Forking returns the ID of the new process. The child process gets its very own process ID. One process calls fork, and if the return is negative one, the fork failed, if it's zero, you're in the child, and if it's anything larger than 0, you have the process ID of the child.

Fork

The fork() system call will spawn a new child process which is an identical process to the parent except that has a new system process ID. The process is copied in memory from the parent and a new process structure is assigned by the kernel. The return value of the function is which discriminates the two threads of execution. A zero is returned by the fork function in the child's process. The environment, resource limits, umask, controlling terminal, current working directory, root directory, signal masks and other process resources are also duplicated from the parent in the forked child process.

```
char * message; int n;
pid_t pid = fork();
switch (pid) {
    case -1:
        perror("fork failed\n");
        exit(1);
    case 0:
        message = "This is the child";
        n = 5;
        break;
    default:
        message = "This is the parent";
        n = 3;
        break; }
for(; n > 0; n--) {
```

```
puts(message);
sleep(1); }
```

Exec

The exec() family of functions will initiate a program from within a program. They are also various front-end functions to execve(). The functions return an integer error code.

The function call execl() initiates a new program in the same environment in which it is operating. An executable (with fully qualified path. i.e. /bin/ls) and arguments are passed to the function. Note that arg0 is the command/file name to execute. Use exec to make a child process execute a new program after it has been forked. It usually does not return, with -1 on failure. It will fail when: too big, ACCESS, get into a loop, the name is too long, or the executable does not exist. The new process keeps the set of blocked signals, pending signals, timers, and any open file descriptors. Doing an exec does not change the relationship between a parent and child process. Caught signals return default values.

```
execl("/bin/ls", "/bin/ls", "-r", "-t", "-l", (
    char *) 0);
```

Wait

wait(): Blocks calling process *until the child process terminates*. If child process has already terminated, the wait() call returns immediately. if the calling process has multiple child processes, the function returns when one returns.

waitpid(): Options available to block calling process for a particular child process not the first one.

Shell scripting

(command >> file appends to file) (command < file content of file to stdin) (|| if left succeeds do not do right)

```
#!/bin/bash
```

```
echo My name is $0
echo My process number is $$
echo I have $# arguments
echo My arguments separately are $*
echo My arguments together are "$@"
echo My 5th argument is '$5'
echo The return value of last command $?
```

```
usage() {
    echo "Usage: $0 filename"
    exit 1
}
is_file_exists(){
    local f="$1"
    [[ -f "$f" ]] && return 0 || return 1
}

# invoke usage
# call usage() function if filename not supplied
[[ $# -eq 0 ]] && usage

# Invoke is_file_exists
if ( is_file_exists "$1" )
```

```

then
    echo "File found"
else
    echo "File not found"
fi

```

Shared memory

Possible for multiple processes to share segments of memory, not full memory. Beneficial for better speed. Access by multiple processes. Still need mutexes and semaphore for synchronizations. Lifetime is independent of any process. All shared memory segments exist in a single list. Identified by a key, generated by the metainfo of files and directories. Processes attach and detach to shared memory segments.

Attach and detach

Used in data structures for shared memory, `shmat` maps shared memory. `shmdt` unmaps shared memory. When using data structures, use shared offsets such as `(struct something) shmat(...)` or use semaphores where second argument isn't zero.

```

if (errno = 0, (key = ftok("/grad/users/morbius",
    42) == -1 )
else if (errno = 0, (shmid = shmget(key, size
    0666, | IPC_REAT | IPC_EXCL)) != -1)
else if (errno = 0, (shmid = shmget(key, 0, 0666)
    ) != -1)

```

Pointers to functions

```

void func(int);
main() {
    void (*fp)(int);
    fp = func;
    (*fp)(1);
    fp(2);
    exit(EXIT_SUCCESS); }
void func(int arg){
    printf("%d\n", arg); }

```

Function pointers provide a way of passing around instructions for how to do something. You can write flexible functions and libraries that allow the programmer to choose behavior by passing function pointers as arguments.

Condition variables

Condition variables provide yet another way for threads to synchronize. While mutexes implement synchronization by controlling thread access to data, condition variables *allow threads to synchronize based upon the actual value of data*.

Without condition variables, the programmer would need to have threads continually polling (possibly in a critical section), to check if the condition is met. This can be very resource consuming since the thread would be continuously busy in this activity. *A condition variable is a way to achieve the same goal without polling*.

A condition variable is *always used in conjunction with a mutex lock*.

Used for waiting. Contains 3 basic operations: wait, signal, broadcast. Wait atomically does unlock and block. Signal atomically does re-locks a thread has it start again.

Deadlock

Four conditions for deadlock: **mutual exclusion** (At least one resource must be held in a non-sharable mode. Only one process can use the resource at any given instant of time.), **circular wait** (A process must be waiting for a resource which is being held by another process, which in turn is waiting for the first process to release the resource.), **no pre-emption** (The operating system must not de-allocate resources once they have been allocated; they must be released by the holding process voluntarily.), **hold and wait** (A process is currently holding at least one resource and requesting additional resources which are being held by other processes.).

Semaphores

A *semaphore* is a special type of variable that can be incremented or decremented, but *crucial access to the variable is guaranteed to be atomic*, even in a multi-threaded program. If two or more threads in a program attempt to change the value of a semaphore, the system guarantees that all the operations will in fact *take place in sequence*.

Contains a counter, not an owner. The counter is anything that is non-negative and can be posted by anything. Wait has two checks: if counter is not zero, decrement and return immediately, but if it is zero then suspend the calling thread and put it in the semaphore queue.

Libraries

Static: Precompiled with all functions. *Advantages:* Simple, reliable. *Disadvantages:* Large, poor security. **Dynamic:** Loaded on the fly. *Advantages:* Typically smaller, good for security. *Disadvantages:* Hard to test and implement.